

# DTrace and Python

Jesús Cea Aviión

jcea@jcea.es

@jcea

<httpS://www.jcea.es/>

<httpS://blog.jcea.es/>



# Jesús Cea Avi3n

- Programming in Python since 1996 (Python 1.4).
- Core Developer since 2008 (Python 2.6 and 3.0).
- Founder of Python Madrid, Python Vigo and Python Espa3a association.
- Solaris user since 1990, SysOP since 1996.
- Consultant and freelance always searching for new and interesting challenges. Hire me! :-)

# Overview of this talk

- What is DTrace?
  - Quick & dirty overview.
- Relevance for Python.
  - Probes in the interpreter.
- Examples tracing a Python program.
- Examples tracing the entire stack, including OS.
- Future. Help!
  - More probes.
  - Porting to other DTrace supported platforms.

# Python

- You already know about this...

# DTrace

- Comprehensive full system dynamic tracing framework developed by Sun Microsystems for Solaris.
- Virtually zero performance impact when not in use.
- Safe to use in production.
- Available on Solaris and derivatives, FreeBSD, NetBSD, Mac OS X, Oracle Linux.

# DTrace

- Operating system, libraries and programs can define “probes”:

```
# dtrace -l | wc -l  
259438
```

- Can fire at machine language function call/return.
- Can fire at arbitrary machine language instruction.
- SAFE to use in production.
- (Almost) zero performance impact when not in use.

# DTrace

- DTrace language is safe, read-only. (\*)
- Probes everywhere:
  - Syscall, virtual memory, CPU scheduler, network, locks, disk...
  - High level probes.
  - Dedicated providers. For instance, Python.
  - Dynamic providers. For instance, sampling profile.
  - Synthetic providers: process defined probes.

# DTrace

- Simple language to activate arbitrary probes and execute code when the event “fires”.
- Speculative tracing.
- It doesn't require process collaboration, but helpful.
- Native aggregation functions.
- Associative arrays.
- Excellent documentation.
- DevOps paradise.

# DTrace examples

Show me the processes doing “fsync()” calls:

```
# dtrace -l -P syscall | wc -l
471 ← Include entry/return + heading
```

```
# dtrace -n 'syscall::fdsync:entry {printf("%s",
execname);}'
dtrace: description 'syscall::fdsync:entry ' matched
1 probe
```

CPU	ID	FUNCTION:NAME
0	58858	fdsync:entry lmtp
[...]		
4	58858	fdsync:entry lmtp
7	58858	fdsync:entry cleanup



# DTrace examples

Show me “fsync()” duration stats:

```
# dtrace -n 'syscall::fdsync:entry {self->t =
timestamp;} syscall::fdsync:return {@t =
quantize(timestamp-self->t);} '
dtrace: description 'syscall::fdsync:entry ' matched
2 probes
^C
```

value	----- Distribution -----	count
262144		0
524288	@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@	7
1048576		0
2097152		0
4194304		0
8388608		0
16777216	@@@@	1
33554432		0

# DTrace examples

Peek inside a process:

```
# dtrace -l -n pid25590:::entry | wc -l  
21204
```

```
# dtrace -l -n pid25590:libssl.so.1.0.0:::entry | wc -l  
649
```

```
# dtrace -n \  
pid9498:libssl.so.1.0.0:ssl_verify_cert_chain:entry  
dtrace: description  
'pid9498:libssl.so.1.0.0:ssl_verify_cert_chain:entry  
' matched 1 probe
```

# DTrace examples

- Sampling profiler:

```
# dtrace -n 'profile-997 /pid == 9912/ {jstack();}'
```

- Show me CPU use of a particular process:

```
# dtrace -n 'BEGIN {oncpu=0;totalcpu=0;} profile-  
997 /pid == 10354/ {oncpu+=1;} profile-997  
{totalcpu+=1;} END {printf("%d %d", totalcpu,  
oncpu);}'
```

```
dtrace: description 'BEGIN ' matched 4 probes
```

```
^C
```

CPU	ID	FUNCTION:NAME
3	2	:END 78972 9871

# DTrace probes in Python

- Instrumented interpreter for better information:

```
# dtrace -l -P python9134
```

ID	PROVIDER	MODULE	FUNCTION	NAME
59421	python9134	libpython3.5m.so.1.0	PyEval_EvalFrameEx	function-entry
59422	python9134	libpython3.5m.so.1.0	PyEval_EvalFrameEx	function-return
59423	python9134	libpython3.5m.so.1.0	_PyGC_CollectNoFail	gc-done
59424	python9134	libpython3.5m.so.1.0	PyGC_Collect	gc-done
59425	python9134	libpython3.5m.so.1.0	gc_collect	gc-done
59426	python9134	libpython3.5m.so.1.0	collect_with_callback	gc-done
59427	python9134	libpython3.5m.so.1.0	_PyGC_CollectNoFail	gc-start
59428	python9134	libpython3.5m.so.1.0	PyGC_Collect	gc-start
59429	python9134	libpython3.5m.so.1.0	gc_collect	gc-start
59430	python9134	libpython3.5m.so.1.0	collect_with_callback	gc-start
59431	python9134	libpython3.5m.so.1.0	subtype_dealloc	instance-delete-done
59432	python9134	libpython3.5m.so.1.0	subtype_dealloc	instance-delete-start
59433	python9134	libpython3.5m.so.1.0	PyType_GenericAlloc	instance-new-done
59434	python9134	libpython3.5m.so.1.0	PyType_GenericAlloc	instance-new-start
59435	python9134	libpython3.5m.so.1.0	PyEval_EvalFrameEx	line

# DTrace probes in Python

- Current probes:
  - line
  - function-entry
  - function-return
  - gc-start
  - gc-done
  - instance-new-start
  - instance-new-done
  - instance-delete-start
  - instance-delete-done

# Examples of DTrace in Python

- Tell me where a particular library call is done:

```
# dtrace -n 'python12042::function-entry
/copyinstr(arg0)=="/usr/local/lib/python3.5/ssl.py
" && copyinstr(a1)=="/getpeercert"/ {jstack(100,
10000);} ' | grep '\[ '
dtrace: description 'python12042::function-entry '
matched 1 probe
[ python3.5/ssl.py:805 (getpeercert) ]
[ urllib3/connection.py:259 (connect) ]
[...]
[ requests/adapters.py:376 (send) ]
[...]
[ requests/api.py:53 (request) ]
[...]
```

# Examples of DTrace in Python

- Tell me how long are garbage collections:

```
# dtrace -n 'python12042:::gc-start {self->t =  
timestamp;} python12042:::gc-done {printf("%uus",  
timestamp-self->t);} '  
dtrace: description 'python12042:::gc-start '  
matched 8 probes  
CPU          ID          FUNCTION:NAME  
   5 101930    collect_with_callback:gc-done 8480us  
   5 101930    collect_with_callback:gc-done 3062us  
   5 101930    collect_with_callback:gc-done 1891us
```

- What Python function fires most GC?
- How frequent are GC?

# Examples of DTrace in Python

- Poor man memory “Leak” detector:

```
# dtrace -n 'python12042:::instance-new-start
{@[copyinstr(arg0)] = sum(1);}
python12042:::instance-delete-done
{@[copyinstr(arg0)] = sum(-1);}'
dtrace: description 'python12042:::instance-new-
start ' matched 2 probes
^C
[...]
```

<code>_GeneratorContextManager</code>	0
<code>socket</code>	0
<code>BufferedSubFile</code>	2
<code>FeedParser</code>	2
<code>HTTPMessage</code>	2



# Examples of DTrace in Python

- Trace a Apache MOD\_WSGI process:

```
# dtrace -n 'python25589:::function-entry
/copyinstr(arg1)== "application" / {self->f=1;}
python25589:::function-return
/copyinstr(arg1)== "application" / {self->f=0;}
python25589:::function-entry /self->f/ {printf("%s",
copyinstr(arg1));}'
```

```
dtrace: matched 3 probes
```

CPU	ID	FUNCTION:NAME
3	4350	PyEval_EvalFrameEx:function-entry application
3	4350	PyEval_EvalFrameEx:function-entry salida

- What operating system calls are being slow?
- Where are we being preempt by the OS? For how long? Why?

# Examples DTrace: Python + OS

- Show me where I am being blocked (synchronization object):

```
# dtrace -n 'sched:::sleep /pid==14857/  
{printf("[BLOCKED] %d\n", tid); jstack();}' | grep  
"\["  
dtrace: matched 7 probes  
  2    5025                cv_block:sleep [START] 2  
[ python3.5/threading.py:293 (wait) ]  
[ python3.5/queue.py:164 (get) ]  
[ python3.5/concurrent/futures/thread.py:64 (_worker) ]  
[...]
```

- CPU accounting per Python Thread.
- What processes are stealing my CPU?
- Examine lock contention, even GIL.

# Examples DTrace: Python + OS

- What code is actually accessing the disk, not getting data from cache?

```
# dtrace -n 'io:::start /pid==14857/ {jstack();}'
```

```
dtrace: description 'io:::start ' matched 6 probes
```

CPU	ID	FUNCTION:NAME
6	5049	bdev_strategy:start
		libc.so.1`_read+0x15
		libpython3.5m.so.1.0`_Py_read+0x4b
		libpython3.5m.so.1.0`_io_FileIO_readall_impl.isra.8+0xeb
		libpython3.5m.so.1.0`PyCFunction_Call+0xca
		libpython3.5m.so.1.0`PyObject_Call+0x68
		libpython3.5m.so.1.0`PyObject_CallMethodObjArgs+0xa2
		libpython3.5m.so.1.0`_io__Buffered_read+0x47f
		libpython3.5m.so.1.0`PyCFunction_Call+0xd9
		libpython3.5m.so.1.0`PyEval_EvalFrameEx+0xa051
		[ <stdin>:1 (<module>) ] ← open("file", "rb").read()
		libpython3.5m.so.1.0`PyEval_EvalCodeWithName+0xb31
		libpython3.5m.so.1.0`PyEval_EvalCode+0x30
		libpython3.5m.so.1.0`PyRun_InteractiveOneObject+0x1a5
		libpython3.5m.so.1.0`PyRun_InteractiveLoopFlags+0x7d
		libpython3.5m.so.1.0`PyRun_AnyFileExFlags+0x40
		libpython3.5m.so.1.0`Py_Main+0xe21
		python3.5`main+0x170
		python3.5`_start+0x80

# Notice:

- You don't modify the source code. You don't even need source code access. No collaboration.
- If you have OS source code, you are GOD!.
- You enable the tracing surgically, when you need it and for the time you need it, from a separate terminal.
- The process continues unaltered, in production.
- Exploratory tracing: hypothesis and fast validation.
- Full system visibility.

# More: Python USDT

*(Userland Statically Defined Tracing)*

- Your python code can define high level probes:
  - client connect, request start, job enqueued, download completed, ...
- Activate logging surgically, on demand, with the daemon running undisturbed.
- You can create individual entry/return probes per function/method with “@fbt” decorator.
- BAD: Stale? code, no documentation. Partial 3.x.

# More: Python USDT

*(Userland Statically Defined Tracing)*

```
Python 2.7.11 (dtrace-issue13405_2.7:8c5948409bbe,  
Mar  3 2016, 04:49:13)  
[GCC 5.3.0] on sunos5  
Type "help", "copyright", "credits" or "license" for  
more information.  
>>> import os  
>>> from usdt.tracer import fbt  
>>> @fbt  
... def example(v) :  
...     pass  
...  
>>> os.getpid()  
24793  
>>> example("hello world!")
```

# More: Python USDT

*(Userland Statically Defined Tracing)*

```
# dtrace -l -P python-fbt24793
```

ID	PROVIDER	MODULE	FUNCTION	NAME
59327	python-fbt24793	fbt	example	entry
59328	python-fbt24793	fbt	example	return

```
# dtrace -n 'python-fbt24793::example:* {}'
```

```
dtrace: description 'python-fbt24793::example:* '
matched 2 probes
```

CPU	ID	FUNCTION:NAME
5	59327	example:entry
5	59328	example:return

# Future:

- Support all DTrace platforms. Sprint tomorrow!
- Add more Python probes in the interpreter and C modules:
  - GIL, Threading module, import machinery...
- Python programs should be able to create personalized dynamic probes. **DONE: Python-USDT.**
- Challenge: integrate with mainstream CPython.



# Performance

- When not enabled, performance hit is **VERY** low:

	DISABLED	ENABLED
0xfee9f79a <+2954>:	jne	0xfee9ede3
0xfee9f7a0 <+2960>:	xor	%eax, %eax
<b>0xfee9f7a2 &lt;+2962&gt;:</b>	<b>nop</b>	<b>int3</b>
0xfee9f7a3 <+2963>:	nop	
0xfee9f7a4 <+2964>:	nop	
0xfee9f7a5 <+2965>:	test	%eax, %eax
0xfee9f7a7 <+2967>:	mov	-0x60(%ebp), %edx
0xfee9f7aa <+2970>:	jne	0xfeea840e

- Current Python USDT implemented in Python, performance hit even when probes are not enabled. Python 2.7, function call+return: x143.

# DTrace Sprint

**tomorrow**

**March, 13<sup>th</sup>**



Main target:

Correctly support FreeBSD,  
NetBSD and Mac OS X.

# Additional References

- Python documentation and code:  
[https://www.jcea.es/artic/python\\_dtrace.htm](https://www.jcea.es/artic/python_dtrace.htm)
- General documentation:  
<https://en.wikipedia.org/wiki/DTrace>  
<http://dtrace.org/guide/preface.html>  
<http://dtrace.org/blogs/>  
<https://wiki.freebsd.org/DTrace/One-Liners>  
[http://dtracebook.com/index.php/Main\\_Page](http://dtracebook.com/index.php/Main_Page)
- Python USDT:  
<https://pypi.python.org/pypi/usdt/>  
<https://github.com/nshalman/python-usdt/>  
<https://github.com/chrisa/libusdt>
- Linux:  
<https://github.com/dtrace4linux/linux>  
[https://docs.oracle.com/cd/E37670\\_01/E37355/html/ol\\_dtrace.html](https://docs.oracle.com/cd/E37670_01/E37355/html/ol_dtrace.html)

# Questions?

- What is “Speculative Tracing”?
- Stack traces and Mac OS X.

<https://www.mail-archive.com/dtrace-discuss@opensolaris.org/msg04668.html>

- What is needed to integrate with mainline CPython?
  - Other interpreters?
- SystemTap synergies.
- DTrace in Linux?

# ¡Thank you!



Jesús Cea Avión  
jcea@jcea.es  
@jcea  
<httpS://www.jcea.es/>  
<httpS://blog.jcea.es/>

